

# **Drilling Deep Into Exadata Performance**

With ASH, SQL Monitoring and Exadata Snapper

**Tanel Põder**  
**Enkitec**

<http://www.enkitec.com>  
<http://blog.tanelpoder.com>

# Intro: About me



ORACLE  
ACE Director

ORACLE | CERTIFIED  
MASTER

- Tanel Põder
  - Former Oracle Database Performance geek
  - Present Exadata Performance geek ;-)
- My Exadata experience
  - I have had the luck to work with all possible Exadata configurations out there
  - Exadata V1 ... X3
  - Multi-rack Exadatas ...
  - Even a mixed rack Exadata (V2 <-> X2-2 :)
- Enkitech Exadata experience
  - Over 100 implementations!



**Expert Oracle Exadata**  
book  
(with Kerry Osborne and  
Randy Johnson of Enkitech)

# About Enkitech



- Enkitech LP
  - North America
- Enkitech Europe
  - EMEA
- 90+ staff
  - In US, Europe
  - Consultants with Oracle experience of 15+ years on average



- What makes us unique
  - **100+ Exadata implementations to date**
    - 100% customer satisfaction rate
  - Exadata-specific services
    - Exadata Quarterly Patching Service
  - Enkitech Exadata lab
    - We have 2 Exadatas for thorough testing and PoCs



## Everything Exadata

Planning/PoC  
Implementation  
Consolidation  
Migration  
Backup/Recovery  
Patching  
Troubleshooting  
Performance  
Capacity  
Training



# Agenda

1. Finding non-Exadata friendly SQL
2. A little bit of Smart Scan internals
3. Reading a SQL Monitoring report on Exadata
  - ... and where it falls short
4. Using advanced Exadata performance metrics
  - Exadata Snapper (ExaSnap)!!!

# Exadata's "secret sauce" for different workloads

- **"DW / Reporting"**

- Long running SQL statements
- Executed less frequently
- Secret Sauce: Smart Scans + Offloading + Storage Indexes
- SQL statements should use full table scans + hash joins

- **"OLTP"**

- Short SQL statements
- Executed very frequently
- Secret Sauce: Flash Cache
- SQL & performance tuning is the same as usual!

## Intro

# Finding non-Exadata-friendly SQL

*(non-smart-scan-friendly SQL...)*

## 1) Finding top non-Exadata-friendly SQLs

- Option 1:
  - Find SQLs which **wait** for non-smart scan IO operations the most
  - **ASH!**
- Option 2:
  - Find SQLs doing **the most disk reads** without Smart Scans
  - Highest MB read or highest IOPS
    - a) ASH!
      - SUM(DELTA\_READ\_IO\_REQUESTS)
      - SUM(DELTA\_READ\_IO\_BYTES)
    - b) ..or join to V\$SQLSTAT (or DBA\_HIST\_SQLSTAT)
      - SUM(PHYSICAL\_READ\_REQUESTS\_DELTA)
      - SUM(PHYSICAL\_READ\_BYTES\_DELTA)

## 2) Are these long-running or frequently executed short queries?

- Exadata Smart scans are not optimized for *ultra*-frequent execution
  - Smart Scanning 1000s of times per second isn't good
  - That's why Smart Scans shouldn't be used for your OLTP-queries
- Here's an analogy:
  1. Want to deliver a single parcel to a destination
    - Use a Ferrari
  2. Want to deliver 10 000 parcels to a destination
    - Use a truck
  3. Want to deliver 10 000 000 parcels to a destination
    - Use a freight train
- **Smart Scan is the Exadata's freight train**
  - Brute force scanning with relatively high inertia to get started, not a few quick (buffered) I/O operations done with surgical precision



## Demo – exafriendly.sql

- Drill down into ASH wait data:


```
SQL> @exadata/exafriendly.sql gv$active_session_history
```

SESSION WAIT_CLASS	EVENT	SECONDS	PCT
ON CPU		192356	34.6
WAITING User I/O	cell single block physical read	191838	34.5
WAITING User I/O	db file parallel read	40577	7.3
WAITING User I/O	cell smart table scan	28593	5.1
WAITING User I/O	cell multiblock physical read	19424	3.5
WAITING User I/O	direct path read temp	18398	3.3
WAITING Application	enq: RO - fast object reuse	8690	1.6
WAITING User I/O	direct path read	5812	1.0
...			

PLAN_LINE	USERNAME	EVENT	SECONDS	PCT
TABLE ACCESS BY LOCAL INDEX ROWID	USER_104	cell single block physical read	40954	21.3
TABLE ACCESS BY INDEX ROWID	USER_779	cell single block physical read	32129	16.7
INDEX RANGE SCAN	USER_104	cell single block physical read	25272	13.2
TABLE ACCESS STORAGE FULL	USER_49	cell single block physical read	9258	4.8
INDEX RANGE SCAN	USER_779	cell single block physical read	4118	2.1
TABLE ACCESS STORAGE FULL	USER_783	cell single block physical read	3641	1.9
UPDATE	USER_104	cell single block physical read	3509	1.8
TABLE ACCESS BY INDEX ROWID	USER_420	cell single block physical read	3341	1.7
MERGE	USER_962	cell single block physical read	2910	1.5

## Demo – mon\_topsql.sql

- TOP Time-consuming SQLs with IO & execution counts
  - That way we'll separate the Ferraris from Freight Trains!
  - The “Ferraris” aren't Exadata *smart scan* friendly



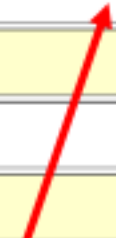
DAY	PCT	OWNER	OBJECT_NAME	PROCEDURE_NAME	SQL_ID	TOTAL_HOURS	TOTAL_SECONDS	EXECUTIONS	SECONDS_PER_EXEC	IO_PCT	CPU_PCT
9/4/11	4.10%				32vklmvdgfkp	40.3	145200	0		65.8	19.1
	3.20%				5qjq8ckgsu054	31.7	114050	0		98.7	1.3
	3.20%				8b1bta0wk04s9	31.7	114010	1159739	0.1	0	98
	3.20%					31.7	113980			2.5	34.1
	3.10%				cwyy1g22pyf60	30.3	109180	0		0	99.5
	2.80%				4suwa3nur4d2q	27.6	99440	0		90.4	6.2
	2.30%				1bnn9jw75t95c	23.1	83130	0		99.3	0.7
	2.10%				gsv18p3ykvv34	20.3	73230	0		0.1	99.9
	1.80%	APPX_MOD_RMS	MAP_PACKAGE	STATE_MAPPINGS	gbmmdsp0ra23n	17.8	63990	7	9141.43	93.3	6
	1.60%				5vgtm5hfhwlwa	15.5	55890	0		98.8	0.8
	1.50%				5xdxcr7pk0cn5	14.7	52810	37985	1.39	98.5	1.5
	1.40%				5fpy9xtn5tv2g	13.6	49130	0		58.8	41.2
	1.20%				3zxcx99ufryt1	12.1	43420	52004462	0	86.9	13.1
9/5/11	2.40%					58.6	211110			7	35.6
	2.40%				fw0tar4rfa39x	58.6	210880	28666	7.36	92.8	5.9
	2.30%				1xxqkv6n9bgu7	58.5	210660	19295	10.92	92.6	6.5
	2.30%				a615cdv2xn65v	58.3	210010	24302	8.64	94	5
	2.30%				fqh4bksfg44bw	58.1	209000	17863	11.7	93.7	5.2
	2.30%				10jbqndnjwuvk	57.5	206980	30286	6.83	86	12.5
	2.00%				cwyy1g22pyf60	49.6	178530	0		0	99.7
	1.90%				4suwa3nur4d2q	47.9	172480	0		94.9	2.4
	1.60%				3zxcx99ufryt1	40	143960	61655359	0	95.6	4.4

## Other sources for overview information

- ASH reports (based on ASH data which I've used)
- EM 12c ASH analytics


### Top SQL with Top Row Sources

SQL ID	PlanHash	Sampled # of Executions	% Activity	Row Source	% RwSrc	Top Event	% Event
<a href="#">91atnkya3uq3u</a>	1846793290	1	6.99	TABLE ACCESS - STORAGE FULL	6.99	cell smart table scan	6.91
<a href="#">a7p9s9nark2aj</a>	256013542	3171	6.89	TABLE ACCESS - STORAGE FULL	6.79	CPU + Wait for CPU	6.79
<a href="#">gctaxcyk0dt67</a>	2150706944	1	2.33	SELECT STATEMENT	2.33	SQL*Net break/reset to client	2.33
<a href="#">7mh3k1p8ht1sy</a>	720331710	1	2.15	TABLE ACCESS - STORAGE FULL FIRST ROWS	2.01	cell multiblock physical read	1.68
<a href="#">1u0xfr31yrh2u</a>	785828209	1	1.76	LOAD AS SELECT	0.60	CPU + Wait for CPU	0.47



# Exadata Architecture

- All DB nodes talk to all (configured) cells - ASM striped data
- **A cell never talks to another cell !!!**



# Smart Scans: Asynchronous, independent prefetching

```
PARSING IN CURSOR #47233445473208 len=38 dep=0 uid=93 oct=3 lid=93 tim=1303079537221426
select * from t3 where owner like 'S%'
END OF STMT
PARSE #47233445473208:c=1000,e=8964,p=0,cr=0,cu=0,mis=1,r=0,dep=0,og=1,plh=4161002650,ti
EXEC #47233445473208:c=0,e=21,p=0,cr=0,cu=0,mis=0,r=0,dep=0,og=1,plh=4161002650,tim=1303
WAIT #47233445473208: nam='SQL*Net message to client' ela= 4 driver id=1413697536 #bytes
WAIT #47233445473208: nam='SQL*Net more data to client' ela= 16 driver id=1413697536 #by
WAIT #47233445473208: nam='reliable message' ela= 1337 channel context=11888341784 chann
WAIT #47233445473208: nam='enq: KO - fast object checkpoint' ela= 143 namelmode=12634685
WAIT #47233445473208: nam='enq: KO - fast object checkpoint' ela= 130 namelmode=12634685
WAIT #47176529789912: nam='cell smart table scan' ela= 25 cellhash#=3176594409 p2=0 p3=0
WAIT #47176529789912: nam='cell smart table scan' ela= 882 cellhash#=2133459483 p2=0 p3=
WAIT #47176529789912: nam='cell smart table scan' ela= 34 cellhash#=3176594409 p2=0 p3=
WAIT #47176529789912: nam='cell smart table scan' ela= 837 cellhash#=2133459483 p2=0 p3=
WAIT #47176529789912: nam='cell smart table scan' ela= 26 cellhash#=2133459483 p2=0 p3=
WAIT #47176529789912: nam='cell smart table scan' ela= 22 cellhash#=379339958 p2=0 p3=0
```

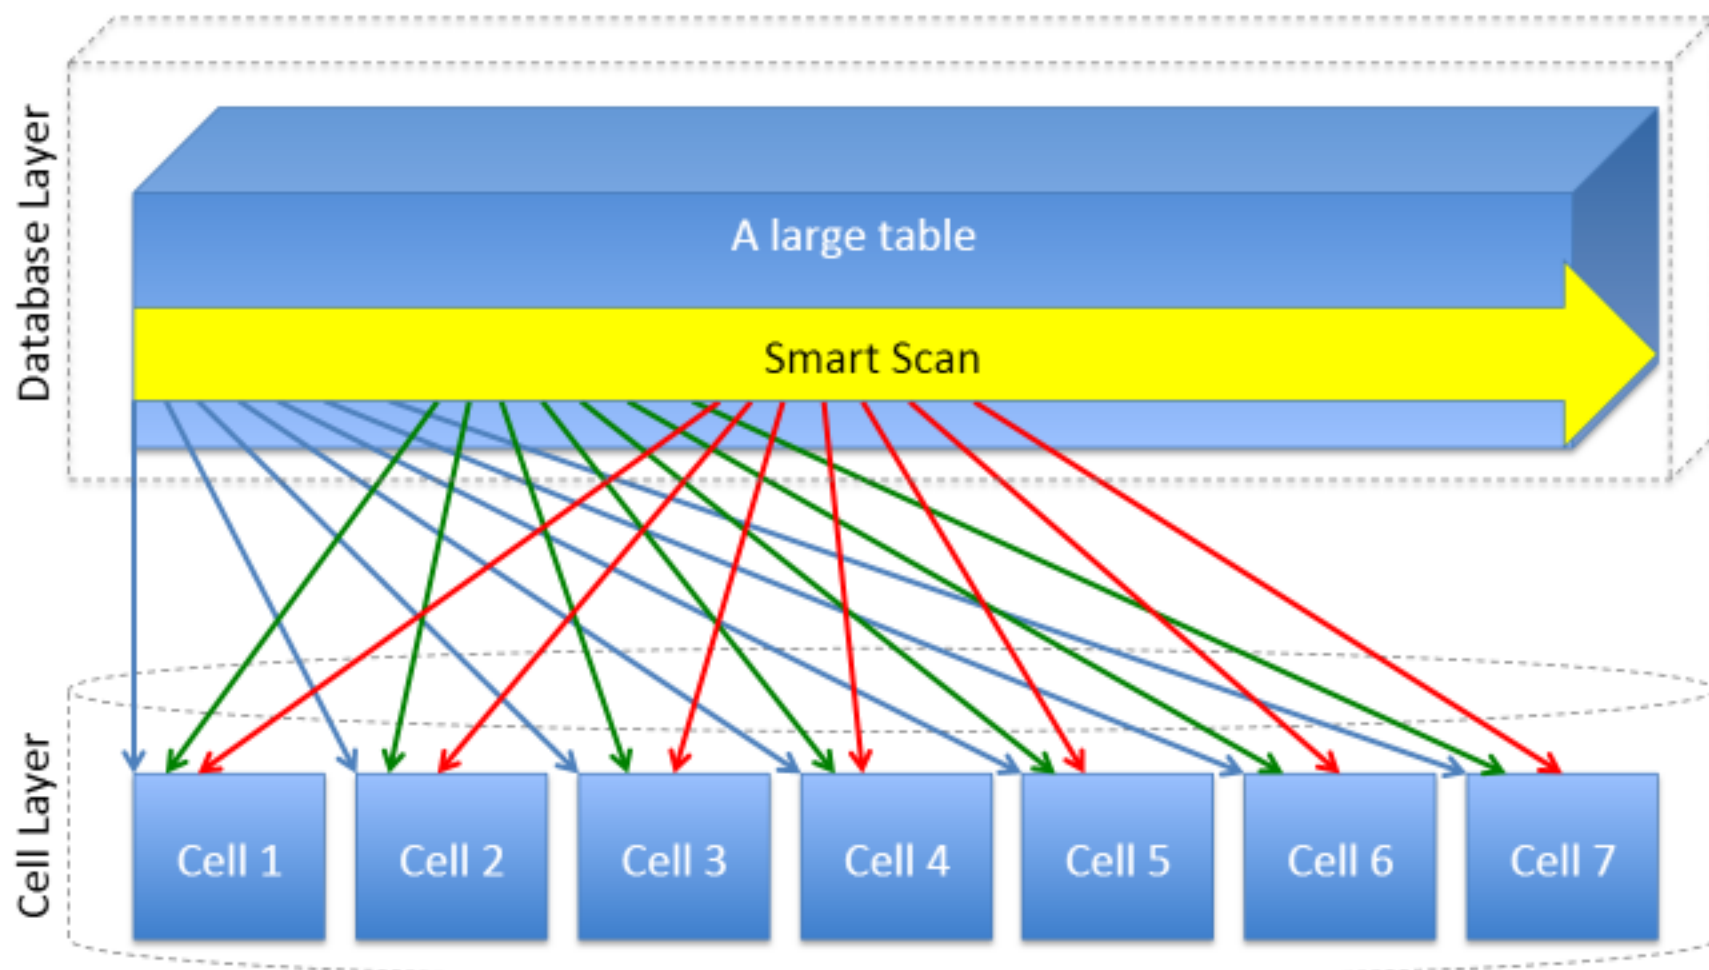
If cell **smart table/index scan** waits show up, then a smart scan is attempted

The waits are so short due to the asynchronous nature of smart scans. Cellsrvs in each cell process the blocks independently to fill their send buffers and the DB just pulls the results from there

The object checkpoint-related wait events **reliable message** and **enq: KO – fast object checkpoint** always precede every direct path scan (thus smart scan too)

## Storage cells are “shared nothing”

- And they don't see what's happening in the database layer...





## Physical disks - Simple math

- Sequential “brute force” scan rate **150 MB/sec** per disk
- or
- **200 random IOPS** per disk
- **12 disks** in a storage cell
- **14 storage cells** in a full rack
- $150 * 12 * 14 = \mathbf{25200 \text{ MB/sec}}$  physical disk scanning rate
  - If doing only sequential brute force scanning
- $200 * 12 * 14 * 8\text{kB} = \mathbf{262.5 \text{ MB/sec}}$ 
  - Random physical read rate (index scan!) around 100x slower!

I'm leaving  
flash cache  
out for  
simplicity for  
now

However,  
Index scans  
can read  
only a subset  
of data

# The Motivation for Writing Exadata Snapper

```
[celladmin@enkcel01 ~]$ cellcli
```

```
CellCLI: Release 11.2.2.4.0 - Production on Mon Mar 05 08:19:41 CST 2012
```

```
Copyright (c) 2007, 2011, Oracle. All rights reserved.
```

```
Cell Efficiency Ratio: 699
```

```
CellCLI>
```

????

## Monitored SQL Execution Details

### Overview

SQL ID 90gw2x39rj2ky ⓘ  
Execution Started Sun Mar 4, 2012 7:07:04 PM  
Last Refresh Time Sun Mar 4, 2012 7:07:10 PM  
Execution ID 16777216  
User TANEL  
Fetch Calls 16

### Time & Wait Statistics

Duration 6.0s  
Database Time 4.2s  
PL/SQL & Java 0.0s  
Wait Activity % 100

### IO Statistics

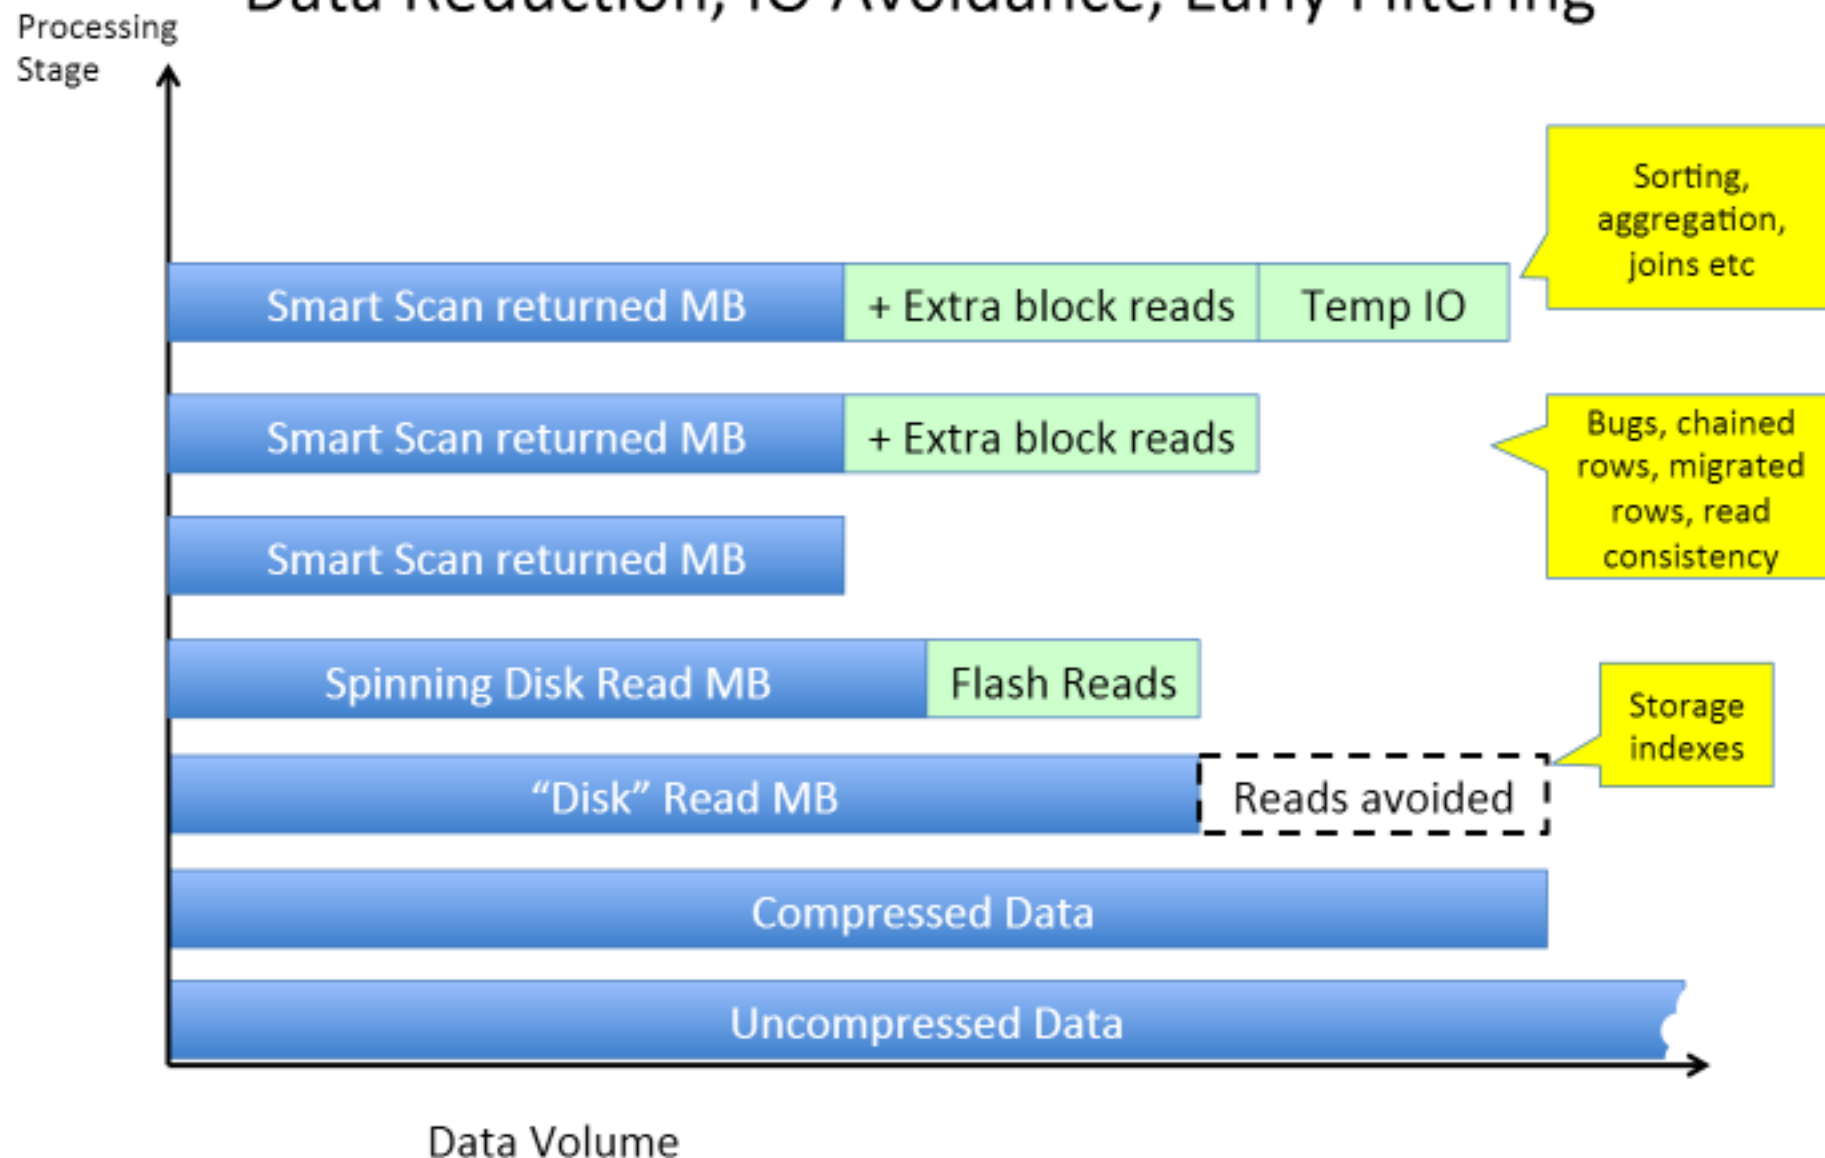
Buffer Gets 1,570  
IO Requests 9,325  
IO Bytes 84MB  
Cell Offload Efficiency -42.86%

Following just the **Cell Efficiency Ratio** can be as misleading as tuning by **Buffer Cache Hit Ratio**!

?



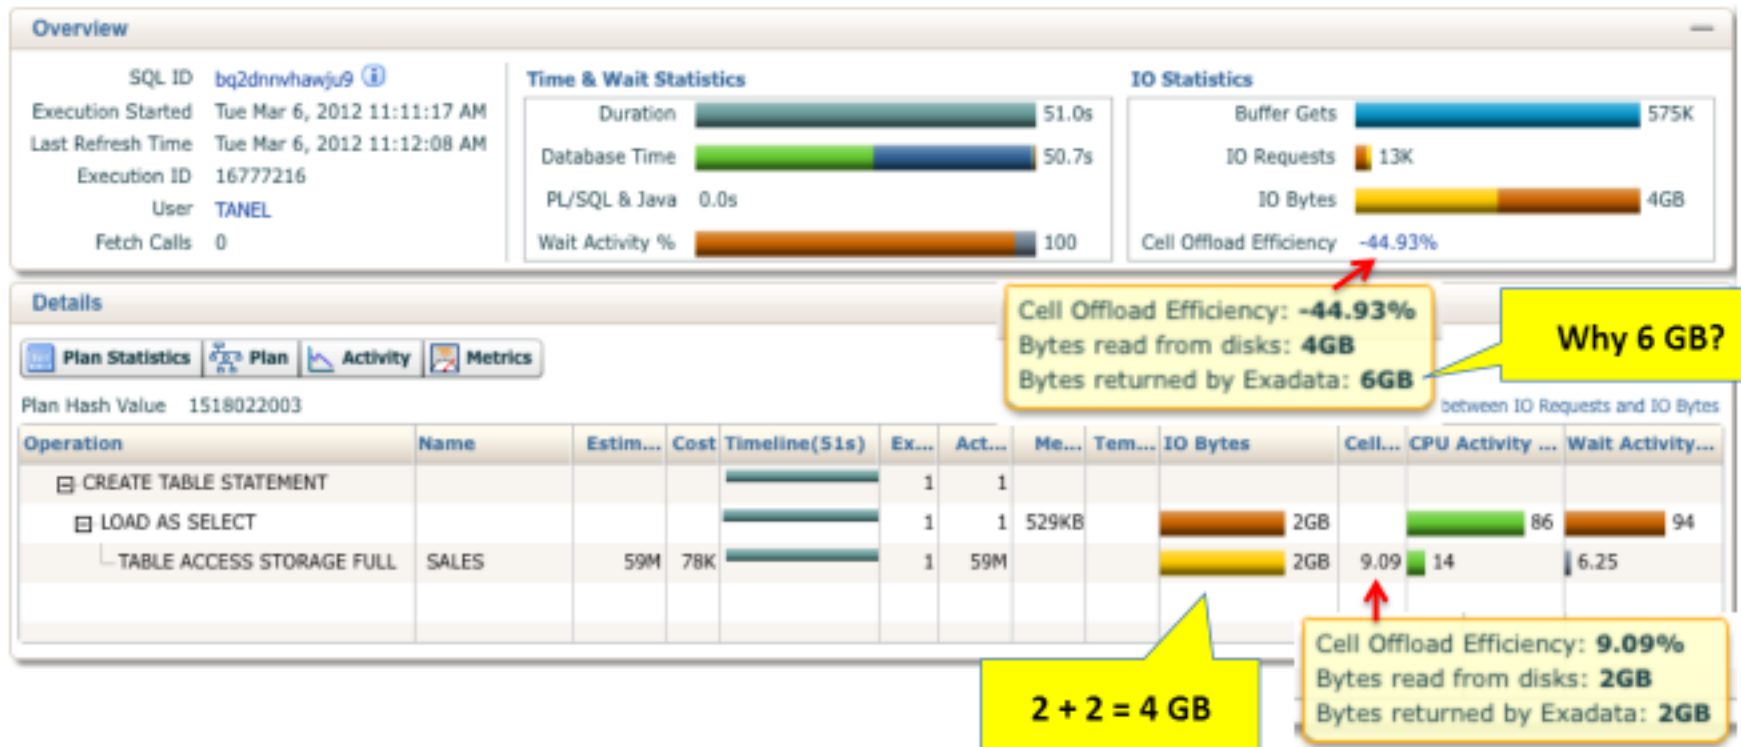
# Data Reduction, IO Avoidance, Early Filtering



## Negative Cell Offload Efficiency ratio?

- Must understand where does this ratio come from
  - Ratio of which Exact metrics?
  - And use those metrics in the future

# Negative Cell Offload Efficiency: Data Load Example



The "Bytes returned by Exadata" metric actually uses the "cell physical IO interconnect bytes" metric internally, which includes all traffic, not just rows returned from smart scan.

So, write IOs also influence cell offload efficiency calculation (data loads, sort/join/aggregate operation TEMP usage).  
**Write IOs are double/triple mirrored by ASM!**

# Interpreting the raw metrics with Exadata Snapper

- IO Efficiency breakdown
  - How much physical IO did we really do?
  - How much reads / how much writes?
  - How much disk IO did we avoid doing thanks to Flash Cache?
  - How much disk IO did we avoid doing thanks to Storage Indexes?
  - What was the total interconnect traffic?
  - How much data was fully processed in the cell (and not shipped back in blocks due to fallback or problems in the cell?)
  - How much data was returned by a Smart Scan (as rows)?
- @exadata/exasnap.sql
  - Beta quality
  - Meant to *complement* ASH and SQL Monitoring reports, not replace them

# ExaSnap example: A CTAS statement

Operation	Name	Estim...	Cost	Timeline(51s)	Ex...	Act...	Me...	Tem...	IO Bytes	Cell...	CPU Activity ...	Wait Activity...
CREATE TABLE STATEMENT					1	1						
LOAD AS SELECT					1	1	529KB		2GB		86	94
TABLE ACCESS STORAGE FULL	SALES	59M	78K		1	59M			2GB	9.09	14	6.25

SQL> @exadata/exasnap % 123 124

-- Exadata Snapper v0.5 BETA by Tanel Poder @ Enkitec - The Exadata Experts (

Cell Offload Efficiency: **-44.93%**  
 Bytes read from disks: **4GB**  
 Bytes returned by Exadata: **6GB**

SID	CATEGORY	METRIC	IOEFF_PERCENTAGE	MB
1280	DB_IO	DB_PHYSIO_MB	#####	4435
	DB_IO	DB_PHYSRD_MB	#####	2219
	DB_IO	DB_PHYSWR_MB	#####	2216
	AVOID_DISK_IO	PHYRD_FLASH_RD_MB		2
	AVOID_DISK_IO	PHYRD_STORIDX_SAVED_MB		0
	DISK_IO	SPIN_DISK_IO_MB	#####	6649
	DISK_IO	SPIN_DISK_RD_MB	#####	2217
	DISK_IO	SPIN_DISK_WR_MB	#####	4432
	REDUCE_INTERCONNECT	PRED_OFFLOAD_MB	#####	2216
	REDUCE_INTERCONNECT	TOTAL_IC_MB	#####	6444
	REDUCE_INTERCONNECT	SMART_SCAN_RET_MB	#####	2009
	REDUCE_INTERCONNECT	NON_SMART_SCAN_MB	#####	4435
	CELL_PROC_DEPTH	CELL_PROC_DATA_MB	#####	2232
	CELL_PROC_DEPTH	CELL_PROC_INDEX_MB		0

The *real* disk writes  
are doubled due to  
ASM double-mirroring

## ExaSnap example 2: Storage Index Savings

The real (spinning) disk reads IO was only **1078 MB** thanks to **1138 MB** of disk IO avoided due to storage indexes:  
(2216 – 1138 = 1078)

```
SQL> @exadata/exasnap basic 90 91
```

```
-- Exadata Snapper v0.5 BETA by Tanel Poder @ Enkitec - The Exadata
```

CATEGORY	METRIC	IOEFF_PERCENTAGE	MB
DB_LAYER_IO	DB_PHYSIO_MB	#####	2216
DB_LAYER_IO	DB_PHYSRD_MB	#####	2216
DB_LAYER_IO	DB_PHYSWR_MB	#####	0
AVOID_DISK_IO	PHYRD_FLASH_RD_MB	#####	0
AVOID_DISK_IO	PHYRD_STORIDX_SAVED_MB	#####	1138
REAL_DISK_IO	SPIN_DISK_IO_MB	#####	1079
REAL_DISK_IO	SPIN_DISK_RD_MB	#####	1078
REAL_DISK_IO	SPIN_DISK_WR_MB	#####	1
REDUCE_INTERCONNECT	PRED_OFFLOADABLE_MB	#####	2216
REDUCE_INTERCONNECT	TOTAL_IC_MB	#####	2
REDUCE_INTERCONNECT	SMART_SCAN_RET_MB	#####	2
REDUCE_INTERCONNECT	NON_SMART_SCAN_MB	#####	0
CELL_PROC_DEPTH	CELL_PROC_DATA_MB	#####	1078
CELL_PROC_DEPTH	CELL_PROC_INDEX_MB	#####	0

All 1078 MB worth of blocks got offloaded: they were opened and processed inside the cells (data layer)

So, why isn't my query Exadata-friendly?

# Drilling down into a SQL execution

## 1. SQL Monitoring report

- Execution plan!
- Where is most of the response time spent (retrieval vs. subsequent processing)
- Are smart scans used for data retrieval?
- IO MB read from disks vs data returned from the cells
  - (also called the offload efficiency ratio but knowing the underlying numbers is way better)


## 2. ExaSnapper

- Or read the raw v\$sesstat metrics if you dare ;-)



# Warm-up case study

- Testing after migration, a query is slow:



## Warm-up case study: check the execution plan

- Parallel execution plan does *not* force full table scans...

Id	Operation	Name	TQ	IN-OUT	PQ Distrib
0	SELECT STATEMENT				
1	PX COORDINATOR				
2	PX SEND QC (RANDOM)	:TQ10002	Q1,02	P->S	QC (RAND)
* 3	FILTER		Q1,02	PCWC	
4	HASH GROUP BY		Q1,02	PCWP	
5	PX RECEIVE		Q1,02	PCWP	
6	PX SEND HASH	:TQ10001	Q1,01	P->P	HASH
7	NESTED LOOPS		Q1,01	PCWP	
8	NESTED LOOPS		Q1,01	PCWP	
9	NESTED LOOPS		Q1,01	PCWP	
10	BUFFER SORT		Q1,01	PCWC	
11	PX RECEIVE		Q1,01	PCWP	
12	PX SEND ROUND-ROBIN	:TQ10000		S->P	RND-ROBIN
* 13	TABLE ACCESS BY GLOBAL INDEX ROWID	ORDERS			
* 14	INDEX RANGE SCAN	ORD_STATUS_IX			
15	PARTITION HASH ITERATOR		Q1,01	PCWP	
* 16	TABLE ACCESS BY LOCAL INDEX ROWID	CUSTOMERS	Q1,01	PCWP	
* 17	INDEX UNIQUE SCAN	CUSTOMERS_PK	Q1,01	PCWP	
18	PARTITION HASH ITERATOR		Q1,01	PCWP	
* 19	INDEX RANGE SCAN	ORDER_ITEMS_PK	Q1,01	PCWP	
20	TABLE ACCESS BY LOCAL INDEX ROWID	ORDER_ITEMS	Q1,01	PCWP	

Serial stage  
in a parallel  
plan

Parallel index  
scans (on  
different  
partitions)

# Warm-up case study – adjusted execution plan

- After forcing full table scans:

Id	Operation	Name	Pstart	Pstop	TQ	IN-OUT	PQ	Di
0	SELECT STATEMENT							
1	PX COORDINATOR							
2	PX SEND QC (RANDOM)	:TQ10003			Q1,03	P->S	QC (RAND)	
* 3	FILTER				Q1,03	PCWC		
4	HASH GROUP BY				Q1,03	PCWP		
5	PX RECEIVE				Q1,03	PCWP		
6	PX SEND HASH	:TQ10002			Q1,02	P->P	HASH	
* 7	HASH JOIN				Q1,02	PCWP		
8	PART JOIN FILTER CREATE	:BF0000			Q1,02	PCWP		
9	PX RECEIVE				Q1,02	PCWP		
10	PX SEND BROADCAST	:TQ10001			Q1,01	P->P	BROADCAST	
* 11	HASH JOIN				Q1,01	PCWP		
12	PART JOIN FILTER CREATE	:BF0001			Q1,01	PCWP		
13	PX RECEIVE				Q1,01	PCWP		
14	PX SEND BROADCAST	:TQ10000			Q1,00	P->P	BROADCAST	
15	PX BLOCK ITERATOR		1	16	Q1,00	PCWC		
* 16	TABLE ACCESS STORAGE FULL	ORDERS	1	16	Q1,00	PCWP		
17	PX BLOCK ITERATOR		:BF0001	:BF0001	Q1,01	PCWC		
* 18	TABLE ACCESS STORAGE FULL	CUSTOMERS	:BF0001	:BF0001	Q1,01	PCWP		
19	PX BLOCK ITERATOR		:BF0000	:BF0000	Q1,02	PCWC		
* 20	TABLE ACCESS STORAGE FULL	ORDER_ITEMS	:BF0000	:BF0000	Q1,02	PCWP		

Fully parallel execution plan

Full table (partition) scans with bloom filter offloading

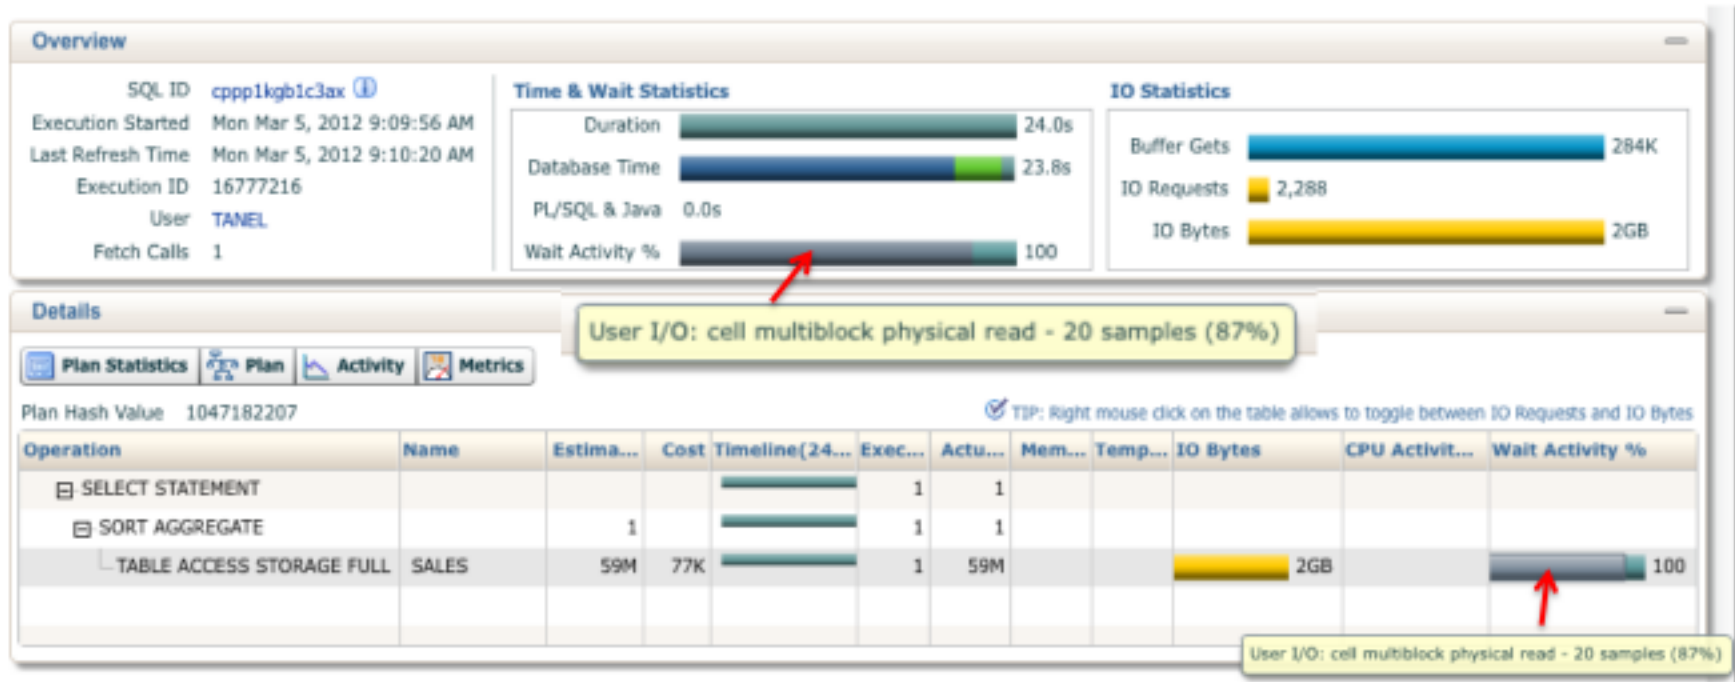
18 - storage(:Z>=:Z AND :Z<=:Z AND ("C"."NLS\_TERRITORY"='AMERICA' AND SYS\_OP\_BLOOM\_FILTER(:BF0000,"C"."CUSTOMER\_ID")))

## Case 2: Response time 24 seconds – where is it spent?

- 1) Full scan: TABLE ACCESS STORAGE FULL
- 2) But waiting for a buffered read wait event

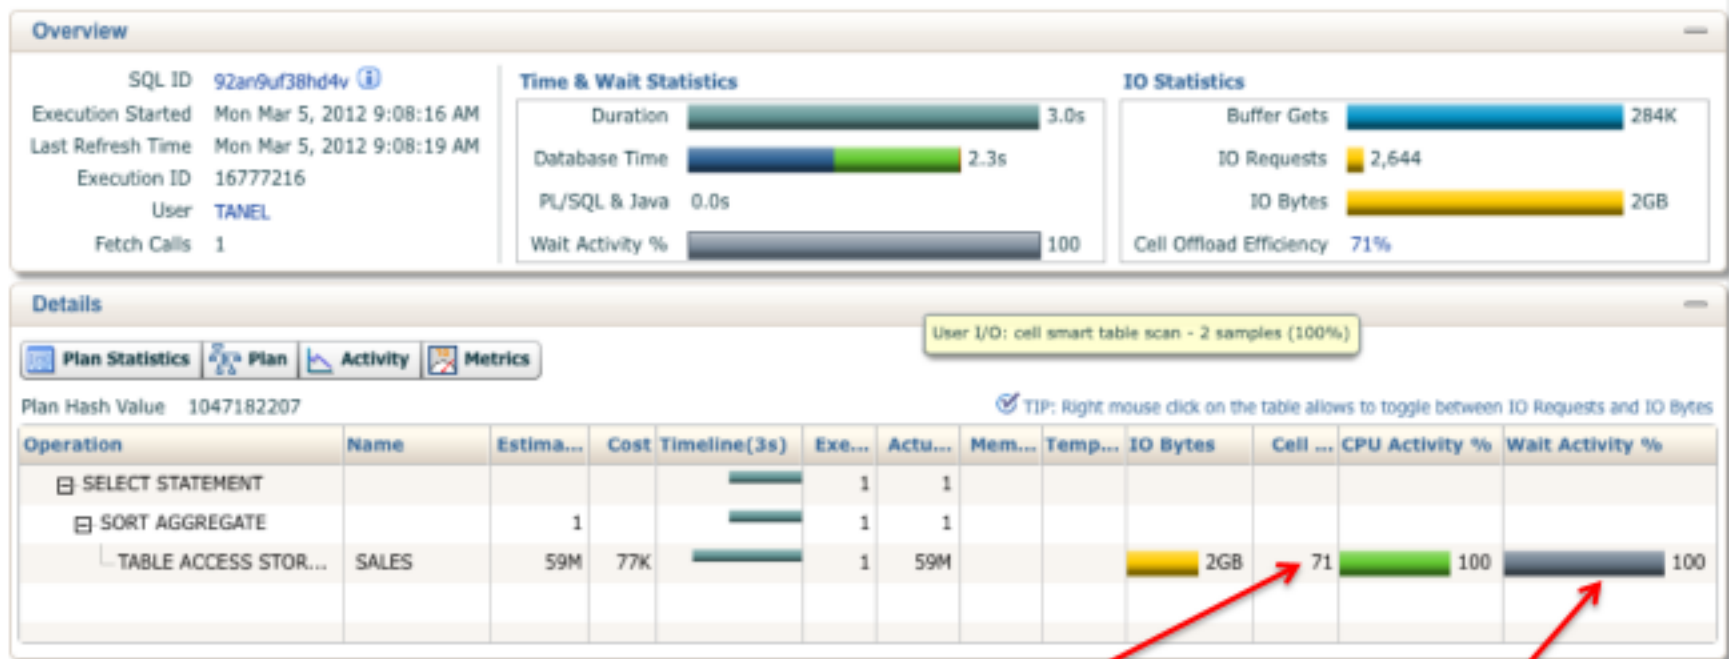
*Smart scan  
will help!*

- *cell multiblock physical read* instead of *cell smart table/index scan*



## Case 2: The same query runs in 3 seconds with Smart Scan

- So, why is it faster?
  - Data retrieval (ACCESS) from storage is faster

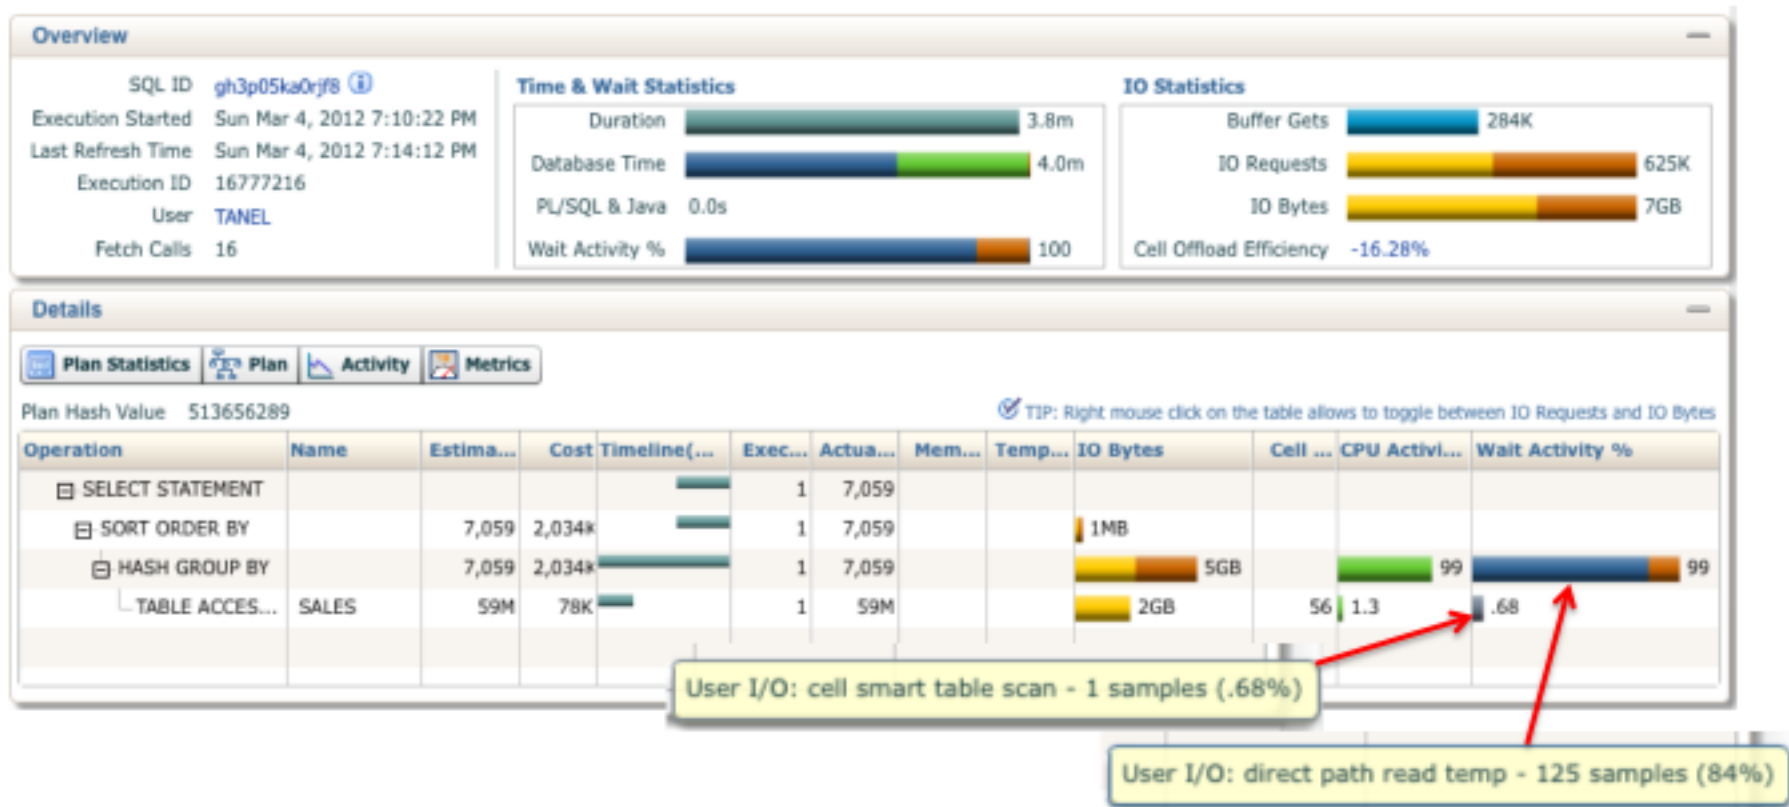


Cell Offload Efficiency: **71%**  
Bytes read from disks: **2GB**  
Bytes returned by Exadata: **648MB**

User I/O: cell smart table scan - 2 samples (100%)

## Case 2: Now let's do something with the data...

- Same query, with an extra GROUP BY:
  - Takes 3.8 minutes! (228 seconds)
  - Why? Let's see what's taking the extra time:



## Checkpoint

- Smart Scans make the **data retrieval from storage** faster
- Any other problems require the usual SQL & DB optimization

Smart Scans do **not** make any of these things faster (by design):

- Index unique/range/skip scans \*
- Sorting
- Aggregation (GROUP BY)
- Analytic Functions
- Filters that can't be pushed down to table level:
  - WHERE **t1.col** + **t2.col** = 1
- Any function calls in projection (select list)
- PL/SQL function calls in projection (or WHERE clause)
- Nested loop joins, sort-merge joins and FILTER lookups
  - Hash joins are special though
- So, you'll have to see where is your bottleneck!
  - SQL Monitoring report is a good start

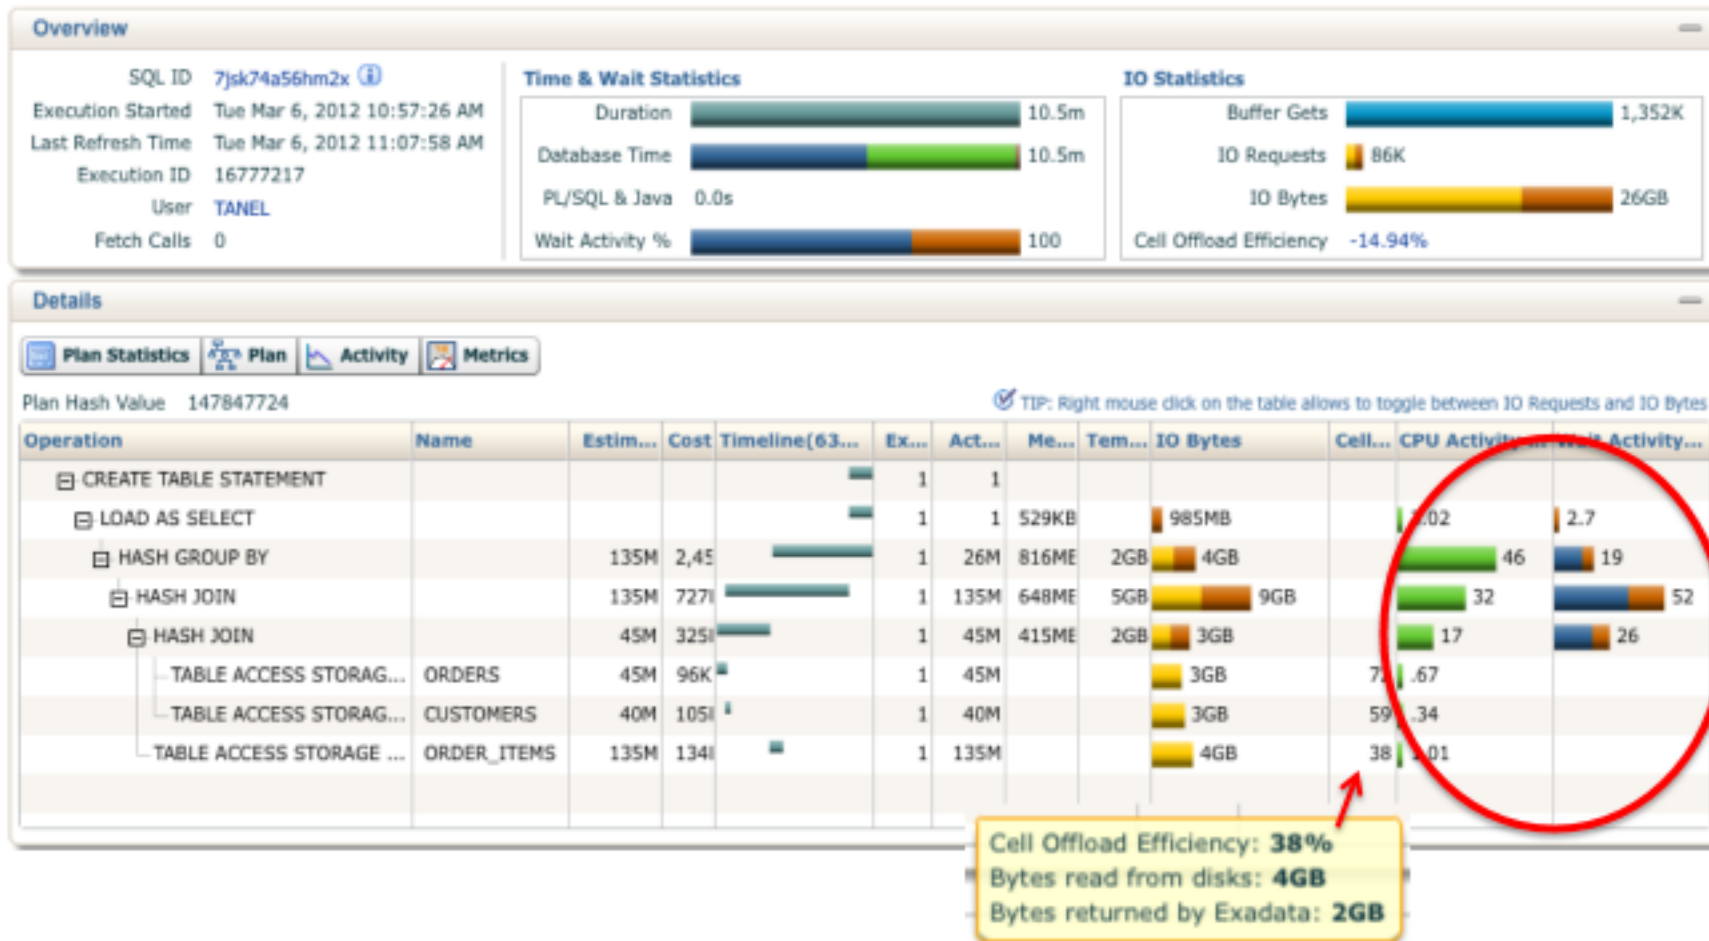
\* A prerequisite for smart scans is a *full segment scan* anyway

This is not a full list of limitations.



# A query bottlenecked by data *processing*, not *retrieval*

- A SQL performing data load and spills to TEMP



## Case 3: Case insensitive search

The **plan** hash value does not change as it doesn't take the predicate differences into account!

```
SELECT SUM(credit_limit) FROM soe.customers  
WHERE cust_first_name LIKE 'j%'
```

Plan hash value: 296924608

FAST ~ 2 seconds

Id	Operation	Name
0	SELECT STATEMENT	
1	SORT AGGREGATE	
* 2	TABLE ACCESS STORAGE FULL	CUSTOMERS

Predicate Information

```
2 - storage("CUST_FIRST_NAME" LIKE 'j%')  
      filter("CUST_FIRST_NAME" LIKE 'j%')
```

```
ALTER SESSION SET nls_comp = LINGUISTIC;  
ALTER SESSION SET nls_sort = BINARY_CI;
```

```
SELECT SUM(credit_limit) FROM soe.customers  
WHERE cust_first_name LIKE 'j%'
```

Plan hash value: 296924608

SLOW ~14 seconds

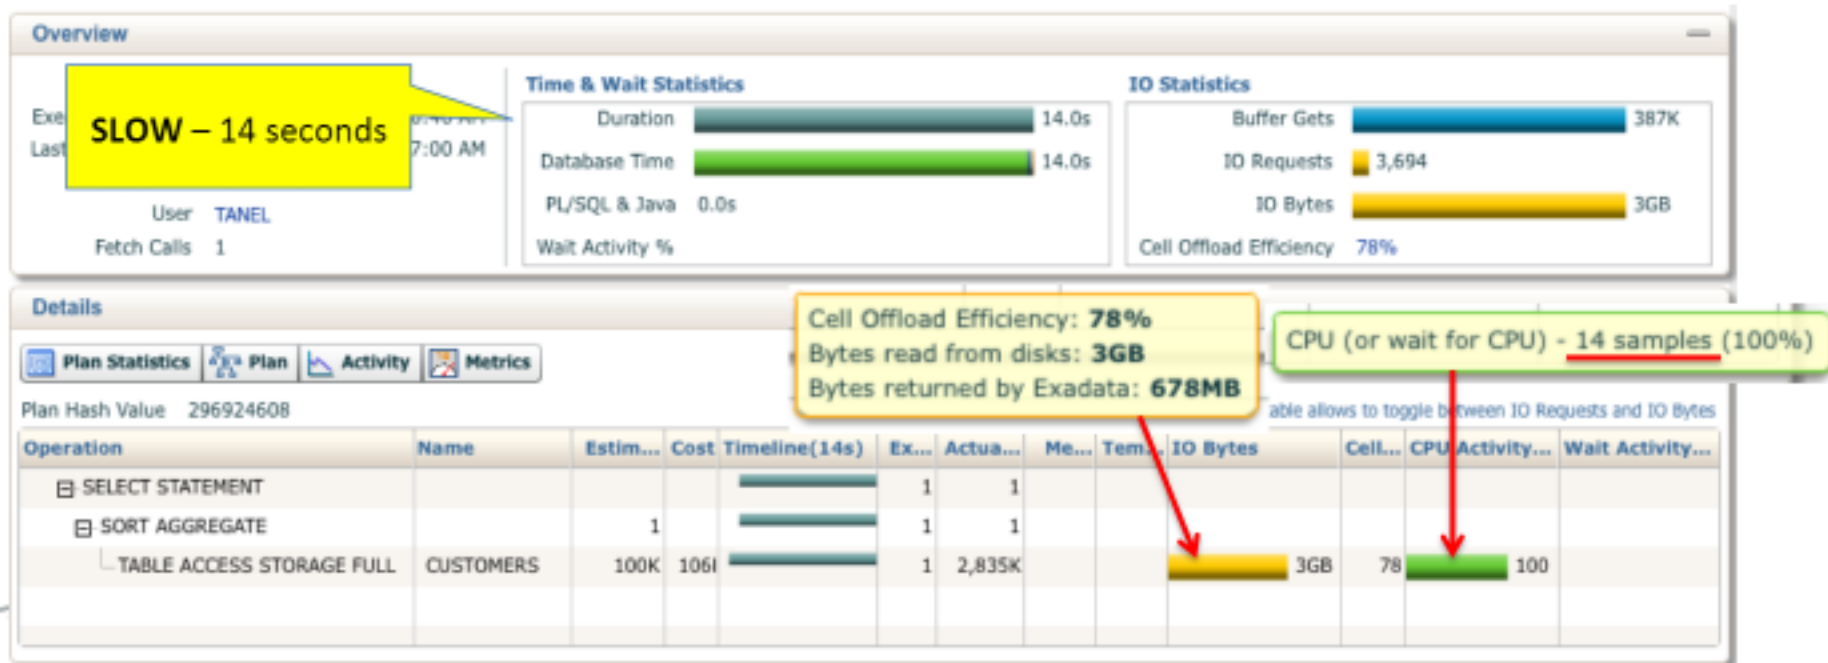
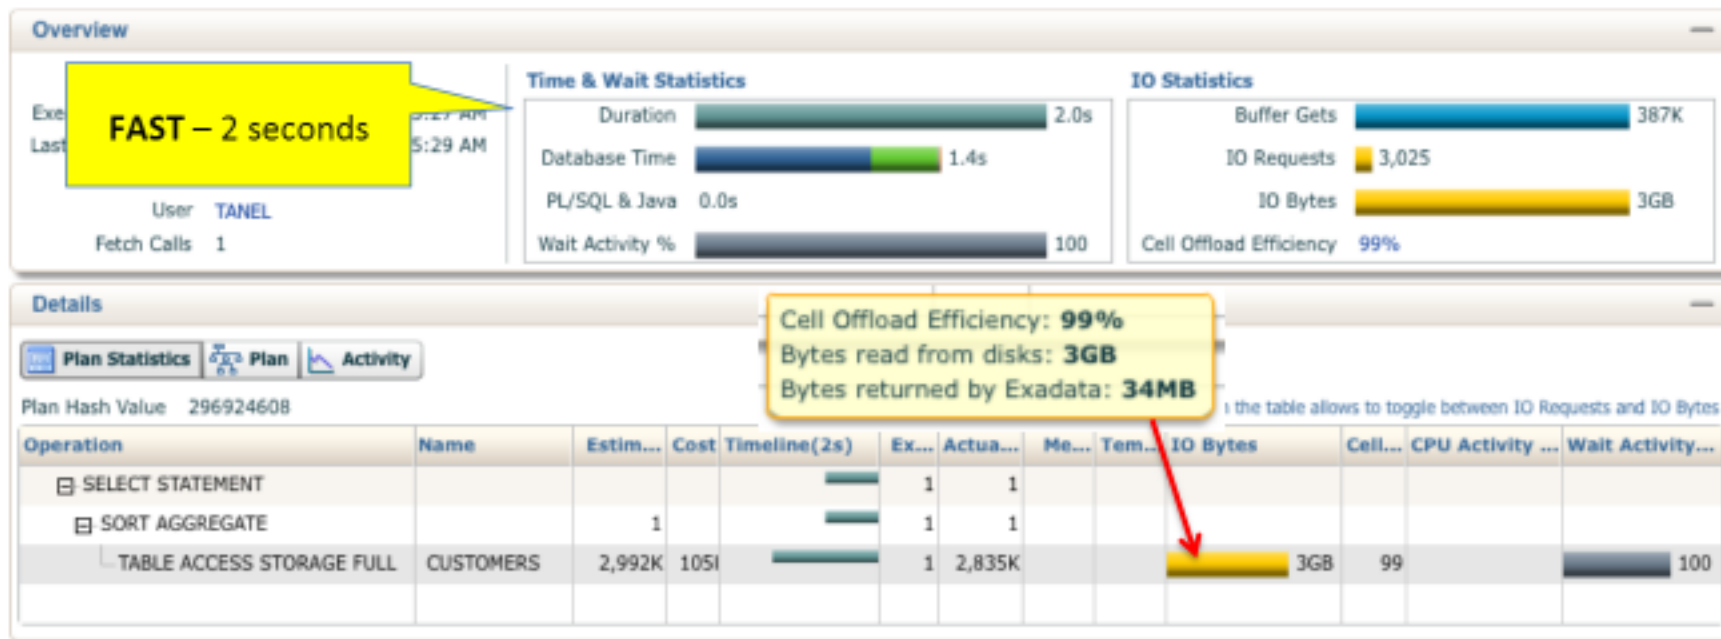
Id	Operation	Name
0	SELECT STATEMENT	
1	SORT AGGREGATE	
* 2	TABLE ACCESS STORAGE FULL	CUSTOMERS

Predicate Information

```
2 - filter(NLSSORT  
      ("OWNER", 'nls_sort=' 'BINARY_CI'')  
      >HEXTORAW('7300'))
```

Where's the  
storage predicate?

[www.enkitech.com](http://www.enkitech.com)



Thank you – and oh, wait!

- **Advanced Exadata Performance seminar!**

- By Tanel Poder
- Systematic Exadata Performance Troubleshooting and Optimization
- 2-day seminar:
  - Dallas, TX – 2-3 May 2013
  - Online – 13-16 May 2013

- <http://blog.tanelpoder.com/seminar/>
- We'll go very deep! ☺

- **Enkitec Extreme Exadata Expo (E4)**

- August 5-6 2013 - lots of great speakers! ->
- <http://enkitec.com/e4>

